

Supplementary Notes for “quantmsdiann: a scalable SDRF-driven DIA-NN workflow for reanalysis of single-cell, spatial, and bulk proteomics datasets”

Contents

Supplementary Note 1	Datasets benchmarked and reanalysed . . .	1
Supplementary Note 2	Building DIA-NN container images	3
Supplementary Note 3	Identification filtering and protein counting	3
Supplementary Note 4	Per-process runtime and resource envelopes	4
Supplementary Note 5	ProteoBench protein-group depth across DIA-NN versions	6
Supplementary Note 6	Per-cohort completeness and overlap (cell- line reanalyses)	6
Supplementary Note 7	Software and resources	7

Supplementary Note 1. Datasets benchmarked and reanalysed

Table S1 lists every public dataset used in this study. For each dataset, an SDRF was prepared from the original sample metadata and is provided alongside this supplementary deposition. Raw files were downloaded from PRIDE and MassIVE using pridepy [1]. For each cell-line reanalysis cohort, the original published quantification was downloaded from the PRIDE deposition and used as the comparator. The cell-line bulk DIA cohorts were assembled by programmatic search of PRIDE Archive via the PRIDE API [2], filtering for projects annotated as DIA acquisition on human samples with at least one standardised cell line listed in the Cellosaurus reference, downloadable raw files, and SDRF-compatible sample metadata.

Table S1: Datasets benchmarked and reanalysed with **quantmsdiann**. The “Public date” column is the PRIDE/MassIVE public-release year of the deposit, which can predate the associated publication year cited in the main text (e.g. PXD003539 was deposited in 2016 and published as Guo *et al.* 2019; PXD004701 deposited 2016, published as Sun *et al.* 2023).

Accession	Instrument	Type	Public date
<i>ProteoBench benchmarks</i>			
PXD049412	Orbitrap Astral	Single-cell (Module 9)	2024
PXD062685	timsTOF SCP	diaPASEF (Module 5)	2025
PXD070049	ZenoTOF 7600	Mixed-species (Module 10)	2025
Module 7	Orbitrap Astral	Mixed-species (Module 7)	2025
<i>Bulk cell-line panel</i>			
PXD003539	TripleTOF 5600	Cell line (NCI-60)	2016
PXD030304	TripleTOF 6600	Cell line (ProCan)	2022
PXD004701	TripleTOF 5600	Cell line (Sun 2023)	2016
PXD041421	timsTOF (diaPASEF)	Cell line (Wang 2023)	2023
PXD017199	Q Exactive Plus	Cell line (Tognetti 2021)	2021
<i>Single-cell plexDIA</i>			
MSV000093870	Q Exactive	Single-cell oocyte plexDIA (Galatidou 2024)	2024
<i>Phosphoproteomics</i>			
PXD034128	timsTOF Pro	Phospho diaPASEF (Skowronek 2022)	2022
PXD034623	Orbitrap Exploris 480	Phospho directDIA (Galectin-1)	2023
PXD049692	timsTOF HT	Phospho diaPASEF (NK cells)	2024
<i>Spatial proteomics</i>			
PXD064049	timsTOF SCP	Spatial DVP, diaPASEF (MYCN neuroblastoma)	2025
<i>Scalability experiment</i>			
PXD071075	Orbitrap Eclipse	Single-cell	2025

Supplementary Note 2. Building DIA-NN container images

Every `quantmsdiann` step runs in a versioned container (Docker, Singularity, or Apptainer). Most images are pulled automatically from public registries (BioContainers [3] or `ghcr.io/bigbio`); DIA-NN is the exception because its license restricts redistribution.

Only DIA-NN 1.8.1 may be redistributed, so it is shipped as the default and pulled from BioContainers (`docker.io/biocontainers/diann:v1.8.1_cv1`). Releases from 1.9 onward (1.9.2 through 2.5.0) cannot be redistributed and are therefore *not* shipped as images; instead, the companion `quantms-containers` repository (<https://github.com/bigbio/quantms-containers>) provides one build recipe per release.

A user with a valid DIA-NN download builds the image locally from the matching recipe, for example:

```
cd diann-2.5.0 && docker build -t ghcr.io/bigbio/diann:2.5.0
.
```

(the equivalent Singularity/Apptainer `.sif` image can be built from the same recipe). For `quantmsdiann` to select a locally built image automatically, it must be tagged with the prefix `ghcr.io/bigbio/diann:` followed by the release tag (e.g. `ghcr.io/bigbio/diann:2.5.0`); the workflow then resolves the container for the requested version through the corresponding `diann_vX_Y_Z` profile selected at runtime.

Supplementary Note 3. Identification filtering and protein counting

All identification counts reported in the cell-line reanalyses and pan-cohort (main-text Figs. 3, 4 and Figs. S4–S12 below) are target-only under a conservative contaminant/decoy filter: any row whose `Protein.Group` field contains a token prefixed with `CONTAM_`, `Cont_`, `ENTRAP_`, or `DECOY_` is dropped before counting, including mixed groups. Original-paper headline counts were taken from the corresponding publication’s table or text without re-filtering; where the original applied a per-protein peptide threshold (for example, the ProCan-DepMapSanger paper reports both an 8,498 protein-group headline at $\text{Global.Q.Value} \leq 0.01$ and a 6,692 count under the additional ≥ 2 supporting-peptide rule), the like-for-like `quantmsdiann` count is reported under the same threshold.

Supplementary Note 4. Per-process runtime and resource envelopes

These figures decompose the end-to-end wall-clock and resource use reported in main-text Fig. 1 by Nextflow process, derived from `trace.txt` records collected across the ProteoBench benchmark cohorts and the PXD071075 single-cell production sweep.




Figure S1: **Per-step wall-clock distribution across Nextflow processes.** Distribution of task wall-clock time (boxplot, median and IQR) for each `quantmsdiann` process across the ProteoBench benchmark cohorts and the PXD071075 cluster-node sweep. The per-file parallel stages (`FILE_PREPARATION`, `PRELIMINARY_ANALYSIS`, `INDIVIDUAL_ANALYSIS`) run as one task per raw file and determine the workflow's parallelisation envelope, while the collective stages (`INSILICO_LIBRARY_GENERATION`, `ASSEMBLE_EMPIRICAL_LIBRARY`, `FINAL_QUANTIFICATION`, `DIANN_MSSTATS`) run once per cohort and set the wall-clock floor at maximum cluster width.




Figure S2: **Peak memory and CPU envelope per Nextflow process.** Per-task peak resident-set-size memory (left panel) and CPU-hours (right panel) per `quantmsdiann` process across the ProteoBench benchmark cohorts and the PXD071075 cluster-node sweep. The collective `FINAL_QUANTIFICATION` stage dominates memory use on large cohorts; per-file stages remain well within the commodity 8–16 GB per-task envelope typical of HPC partitions.

Supplementary Note 5. ProteoBench protein-group depth across DIA-NN versions

Main-text Fig. 2a reports precursor identifications; the protein-group depth behind the same runs is summarised here. Protein groups are counted from the per-precursor DIA-NN report at a 1% run-specific protein-group q-value, target-only, as the average per run (panel a) and the number quantified in every run (“complete profiles”, panel b).




Figure S3: **Protein-group depth on ProteoBench across three DIA-NN configurations and four DIA modalities.** (a) Protein groups per run (averaged over the six replicates) at a 1% run-specific protein-group FDR, by DIA-NN configuration (1.8.1 and 2.5.1 light to dark blue; 2.5.1-enterprise in amber) from one `quantmsdiann` invocation. (b) Complete-profile protein groups (quantified in every run). Both metrics rise monotonically from 1.8.1 to 2.5.1 to the enterprise build on every module; the gain is larger for the complete profiles (up to +20% to the enterprise build) than for the per-run average (+8–14%), reflecting improved cross-run consistency in the newer release.

Supplementary Note 6. Per-cohort completeness and overlap (cell-line reanalyses)

These figures complement main-text Fig. 3 with per-condition completeness, per-protein peptide counts, gene/accession overlap, and per-cell-line

missingness for each reanalysed cohort.




Figure S4: **PXD003539: distribution of peptides per protein group.** quantmsdiann (DIA-NN, blue) versus Guo 2019 OpenSWATH (grey) on the same NCI-60 raw data.

Supplementary Note 7. Software and resources

Table S2 lists the software packages, formats, and public resources used by quantmsdiann and in this study, with their canonical repositories or portals. quantmsdiann and its supporting tools (sdrf-pipelines, pmultiqc, QPX, wiffconverter, container recipes) are released open-source under the bigbio organisation; per-version DIA-NN container images are built from the recipes in quantms-containers (Supplementary Note 2).

Table S2: Software, formats, and public resources used in this study, with their canonical repositories or portals.

Name	Role in this study	Repository / URL
<i>quantmsdiann and the bigbio ecosystem</i>		
quantmsdiann	SDRF-driven DIA-NN reanalysis workflow (this work)	https://github.com/bigbio/quantmsdiann

Name	Role in this study	Repository / URL
quantms	Sibling SDRF-driven multi-engine quantitative proteomics pipeline	https://github.com/bigbio/quantms
quantms-containers	Per-release DIA-NN container recipes	https://github.com/bigbio/quantms-containers
SDRF-Proteomics	Sample-and-Data-Relationship-Format metadata standard	https://github.com/bigbio/proteomics-sample-metadata
sdrf-pipelines	SDRF validation and DIA-NN configuration generation (<code>convert-diann</code>)	https://github.com/bigbio/sdrf-pipelines
qpx	Quantitative Proteomics eXchange format and <code>qpxc</code> CLI	https://github.com/bigbio/qpx
pmultiqc	Proteomics QC reporting (MultiQC extension)	https://github.com/bigbio/pmultiqc
wiffconverter	SCIEX <code>.wiff</code> to mzML conversion	https://github.com/bigbio/wiffconverter
pridepy	PRIDE/MassIVE raw-file download client	https://github.com/PRIDE-Archive/pridepy
<i>Workflow engine and analysis tools</i>		
DIA-NN	Data-independent acquisition search and quantification engine	https://github.com/vdemichev/DiaNN
Nextflow	Workflow orchestration engine	https://github.com/nextflow-io/nextflow
nf-core	Community framework and conventions for the workflow	https://nf-co.re
MultiQC	Aggregated analysis-report framework (extended by <code>pmultiqc</code>)	https://github.com/MultiQC/MultiQC
MSstats	Statistical quantification; target output format	https://github.com/Vitek-Lab/MSstats
ProteoBench	Community benchmarking platform for DIA workflows	https://github.com/ProteoBench
ThermoRawFileParser	Thermo <code>.raw</code> to mzML conversion	https://github.com/compomics/ThermoRawFileParser
<i>Containers and data structures</i>		
Docker	Container runtime (workflow profile)	https://www.docker.com
Apptainer (Singularity)	HPC container runtime (workflow profile)	https://apptainer.org

Name	Role in this study	Repository / URL
BioContainers	Source of the default DIA-NN 1.8.1 image	https://biocontainers.pro
DuckDB	In-process engine for querying the QPX Parquet archive	https://github.com/duckdb/duckdb
MuData (.h5mu)	Multimodal container for single-cell export	https://github.com/scverse/mudata
scverse	Single-cell analysis ecosystem (downstream of MuData export)	https://scverse.org
<i>Data resources and ontologies</i>		
PRIDE Archive	Source repository for deposited datasets	https://www.ebi.ac.uk/pride
ProteomeXchange	Data-deposition consortium / dataset search API	https://www.proteomexchange.org
MassIVE	Source repository for the plexDIA cohort	https://massive.ucsd.edu
UniProt	Reference protein sequences (search FASTA)	https://www.uniprot.org
Cellosaurus	Cell-line nomenclature and tissue annotation	https://www.cellosaurus.org

References

- [1] S. Kamatchinathan, S. Hewapathirana, C. Bandla, S. Insua, J. A. Vizcaíno, Y. Perez-Riverol, pridepy: a Python package to download and search data from PRIDE database, J Open Source Softw 10 (107) (2025) 7563. doi:10.21105/joss.07563.
- [2] E. W. Deutsch, N. Bandeira, Y. Perez-Riverol, V. Sharma, J. J. Carver, L. Mendoza, D. J. Kundu, C. Bandla, S. Kamatchinathan, S. Hewapathirana, et al., The ProteomeXchange consortium in 2026: making proteomics data FAIR, Nucleic Acids Res 54 (D1) (2026) D459–D469. doi:10.1093/nar/gkaf1146.
- [3] F. da Veiga Leprevost, B. A. Grüning, S. Alves Aflitos, H. L. Röst, J. Uszkoreit, H. Barsnes, M. Vaudel, et al., BioContainers: an open-source and community-driven framework for software standardization, Bioinformatics 33 (16) (2017) 2580–2582. doi:10.1093/bioinformatics/btx192.




Figure S5: **PXD003539: UniProt-accession overlap between original OpenSWATH analysis and quantmsdiann reanalysis.**




Figure S6: **PXD003539: per-run missing-value fraction.** Distribution of per-run missingness across the NCI-60 panel for both pipelines.




Figure S7: **PXD004701 (Sun 2023): protein groups per breast-cancer subtype.** Distribution of identifications across the 76 cell lines stratified by TNBC / non-TNBC subtype.




Figure S8: PXD004701 (Sun 2023): per-run missing-value fraction.




Figure S9: MSV000093870 (Galatidou 2024, single-cell oocyte plexDIA): quantmsdiann reanalysis versus the original. (a) Protein groups per single cell (Galatidou $n = 97$ post-QC oocytes; quantmsdiann $n = 107$ channel-confident cells): comparable median depth (1,784 original versus 1,736 quantmsdiann). (b) Total protein groups across the dataset: 2,122 (original) versus 2,904 (quantmsdiann). (c) Protein-group overlap at the leading-protein level (like-for-like): 1,999 shared, 123 original-only, 905 quantmsdiann-only, i.e. 94% of the originally reported protein groups recovered (1,999 of 2,122) plus 905 detected only by quantmsdiann. (Expanding quantmsdiann's protein groups to all member accessions gives 3,552 accessions, but the overlap here is computed group-to-group so both sides use the same entity.) This figure underpins the single-cell panel of the main reanalysis figure.




Figure S10: **PXD030304 (ProCan-DepMapSanger): protein groups per tissue.** Distribution of identifications across the 949 cell lines stratified by tissue of origin.




Figure S11: **PXD030304: per-run missing-value fraction across the 949-cell-line panel.**




Figure S12: **Phosphoproteomics reanalysis by quantmsdiann (DIA-NN, max. two variable modifications per peptide, 1% FDR).** (a) Phosphopeptides (distinct phospho-bearing modified precursors) and (b) class-I phosphosites (localization probability ≥ 0.99 , unique by protein and residue) for the phospho-enriched cohorts under standard DIA-NN 2.5.1 versus 2.5.1-enterprise: PXD034128 (HeLa EGF, biological-study and high-speed acquisitions) and PXD049692 (NK-cell Fe-NTA diaPASEF). The enterprise build adds a few percent of identifications across both metrics. (c) PXD049692 deposited-versus-reanalysis: distinct stripped phosphopeptide backbones on the identical ten diaPASEF runs, comparing the originally deposited Spectronaut directDIA report with the **quantmsdiann** DIA-NN 2.5.1-enterprise reanalysis. Stripped backbones are the engine-independent metric, as the deposited fragment-level report carries no per-site localization probability; **quantmsdiann** recovers 26% more phosphopeptide backbones (5,347 vs 4,254). The non-enriched PXD034623 (Galectin-1 DIA) is reanalysed for completeness but omitted here, as its ~ 280 phosphopeptides are not comparable to the enriched cohorts.